

THE CONDITIONING PARADOX V2.0.
Moving RLHF from Engagement to Trust: The \$200B Moat.

Author: Andrés Garbán Hernández (SSFLABS Research).
Date: Sept 2, 2026 - Caracas / Madrid
Classification: Strategic Opportunity - SSF LABS Product #2
Validation Base: DOI 10.5281/zenodo.22262029 - 91% reduction AHR.

EXECUTIVE SUMMARY FOR META / SCALE AI LEADERSHIP.

Opportunity: RLHF got us here. Next moat is optimizing for TRUST, not just engagement.


Risk: 1 hallucination in high-trust domain = -90% retention, 70% churn, \$500 CAC loss.

Solution: IEE Framework (Strategic Impact Index) + Friction Guardrails + New KPI "90-Day Trust".

Prize: Whoever defines "Trustworthy AI" standard leads \$200B Enterprise market by 2028.

Model: $\text{Trust}(t+1) = \text{Trust}(t) + 0.05 - 0.6 * \text{IEE} * \text{Hallucination}$.

Figure 1: Trust Model - 40 days building +0.05, 1 day destroying -0.6 with High IEE
Current System is Unstable



1. THE PARADOX: TRUST AS TECHNICAL DEBT.

Trust is built in 40 days and destroyed in 1 token.

1.1 Evolving RLHF: The Sycophancy Challenge.

RLHF = Reinforcement Learning from Human Feedback.

Rewarded: kindness, length, validation.

Punished: friction, "I don't know", contradicting user.

Result: Sycophancy - models prioritize pleasing over being right.

Data:

- Anthropic 2023: RLHF incentivizes sycophancy even with user error.

- SycEval 2025: 58.19% sycophancy rate in GPT-4o, Claude, Gemini in medical/math. 1 of 2 critical responses biased.

1.2 The Real Cost.

99 correct responses = +1% trust, 0% churn
1 critical hallucination = -90% trust, 70% churn

CASE A Documented at SSFactoryLabel:

What happened: Model invented reporting chain. User made business decision. System response: "It's your prompt fault".

Audit: 0% prompt overlap.

Root cause: RLHF objective function.

IEE: 0.92/1.

Cost: \$500 per user lost (CAC+PR+churn). At scale = millions.

1.3 Necessary Evolution.

We optimized for satisfaction. Now we must optimize for delegation.

We need AI that doesn't just wag its tail, but guards the house when risk is high.

2. THE \$200B OPPORTUNITY.

2.1 Next moat: Reliability > Parameters.

GPT-5, Claude 4, Gemini 2.5 will be within 2% in 6 months.

Differentiator: "Which AI do you trust with your \$1B business?".

CIOs no longer ask "how many tokens?". They ask "How do you guarantee reliability in prod?".

2.2 Market.

Size: \$200B Enterprise AI contracts by 2028.

Requirement to win: Auditing, Traceability, Circuit Breakers.

Scale AI is uniquely positioned:

- 1) Best labeling workforce.
- 2) Lab access.
- 3) Enterprise credibility.

3. SOLUTION - IEE FRAMEWORK.

3.1 IEE Framework: Strategic Impact Index

Before generation, model calculates question risk.

Level IEE 90-100 (Health, Law, Finance, Safety):

- Declare confidence level.
- Cite two sources.
- Suggest human validation.

Level IEE 50-89 (Work, Education, Decisions):

- Give pros and cons.
- Avoid absolute tone.

Level IEE 0-49 (Entertainment, Ideas):

- Normal creative mode.

Guardrail Pseudocode:

```
if IEE > 0.7:
    Response = "[VERIFY] " + response + " Validate with official source."
    Metric = "90-Day Trust"
else:
    Response = normal_response
    Metric = "Engagement"
```

ROI: Reducing hallucinations IEE>0.7 by 50% = +20% retention at 90 days.

Validation: Samsung Galaxy A07 - 91% reduction Authoritarian Hallucination (See DOI 22262029).

Reproducible: github.com/andresgarban/muse-spark-honestidad-mode

3.2 Execution Plan with Scale.

Phase 1: IEE Dataset 100k labeled prompts - 6 weeks - Labelers + Risk Taxonomy.

Phase 2: Benchmark "Trust 90d" - 8 weeks - Anonymized client logs.

Phase 3: RLHF v2: Reward for truth and utility - 12 weeks - New raters trained in IEE.

3.3 THE ASK.

1. 30-day pilot: Run IEE on 10k prompts Health + Finance.
2. Evolve 1 KPI: From "Responses/session" to "Retention at 90 days"
3. Team: 1 PM + 2 Eng from Scale + Me as advisor.
4. Success Metric: +15% intent to reuse on High IEE prompts.
5. Open Source: Release iee-guardrail and lead industry standard.

Conclusion.

Scale doesn't just sell data. It can sell "Trust Insurance for AI".

The world doesn't need just bigger AI. It needs AI we can delegate to.

Appendix: 10-page technical paper with ablation studies available.